

Automated, Retargetable Back-Annotation for Host Compiled Performance and Power Modeling

Zhuoran Zhao, Suhas Chakravarty and Andreas Gerstlauer
Electrical and Computer Engineering, The University of Texas at Austin

Introduction

Motivation

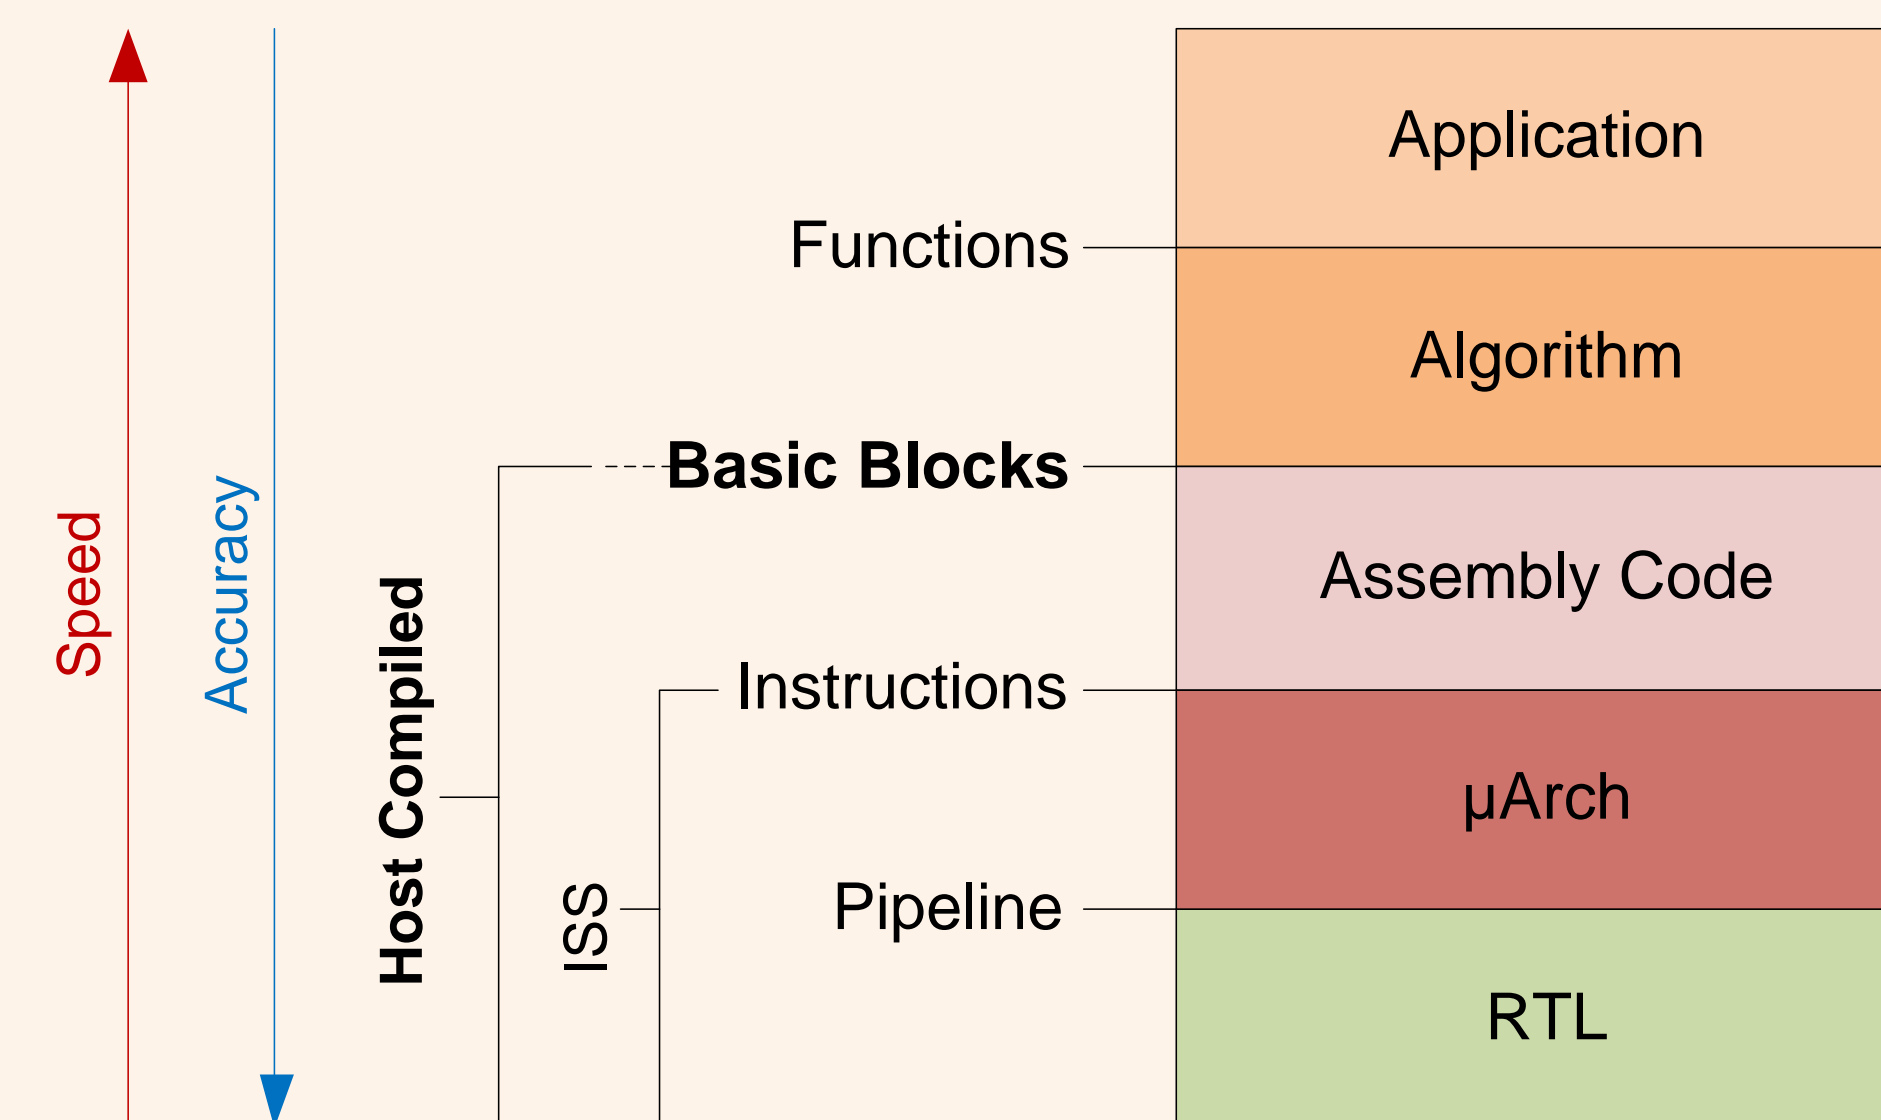
- Increasing design complexities
 - Rapid design space exploration desired
 - Performance and power validation
- Traditional simulation models
 - Instruction Set Simulator (ISS)
 - RTL/Gate level
 - Too slow or too inaccurate
- Modeling at higher abstraction levels
 - Fast and accurate
 - Host-compiled simulation

Challenges & Solutions

- Annotation granularity?
 - Speed vs. accuracy tradeoff
 - Block (BB) granularity
- Compiler optimizations?
 - Mapping between source and binary
 - Work with intermediate representation (IR)
- Dynamic architecture effects?
 - Pipelining state
 - Pairwise characterization

Host-Compiled Modeling

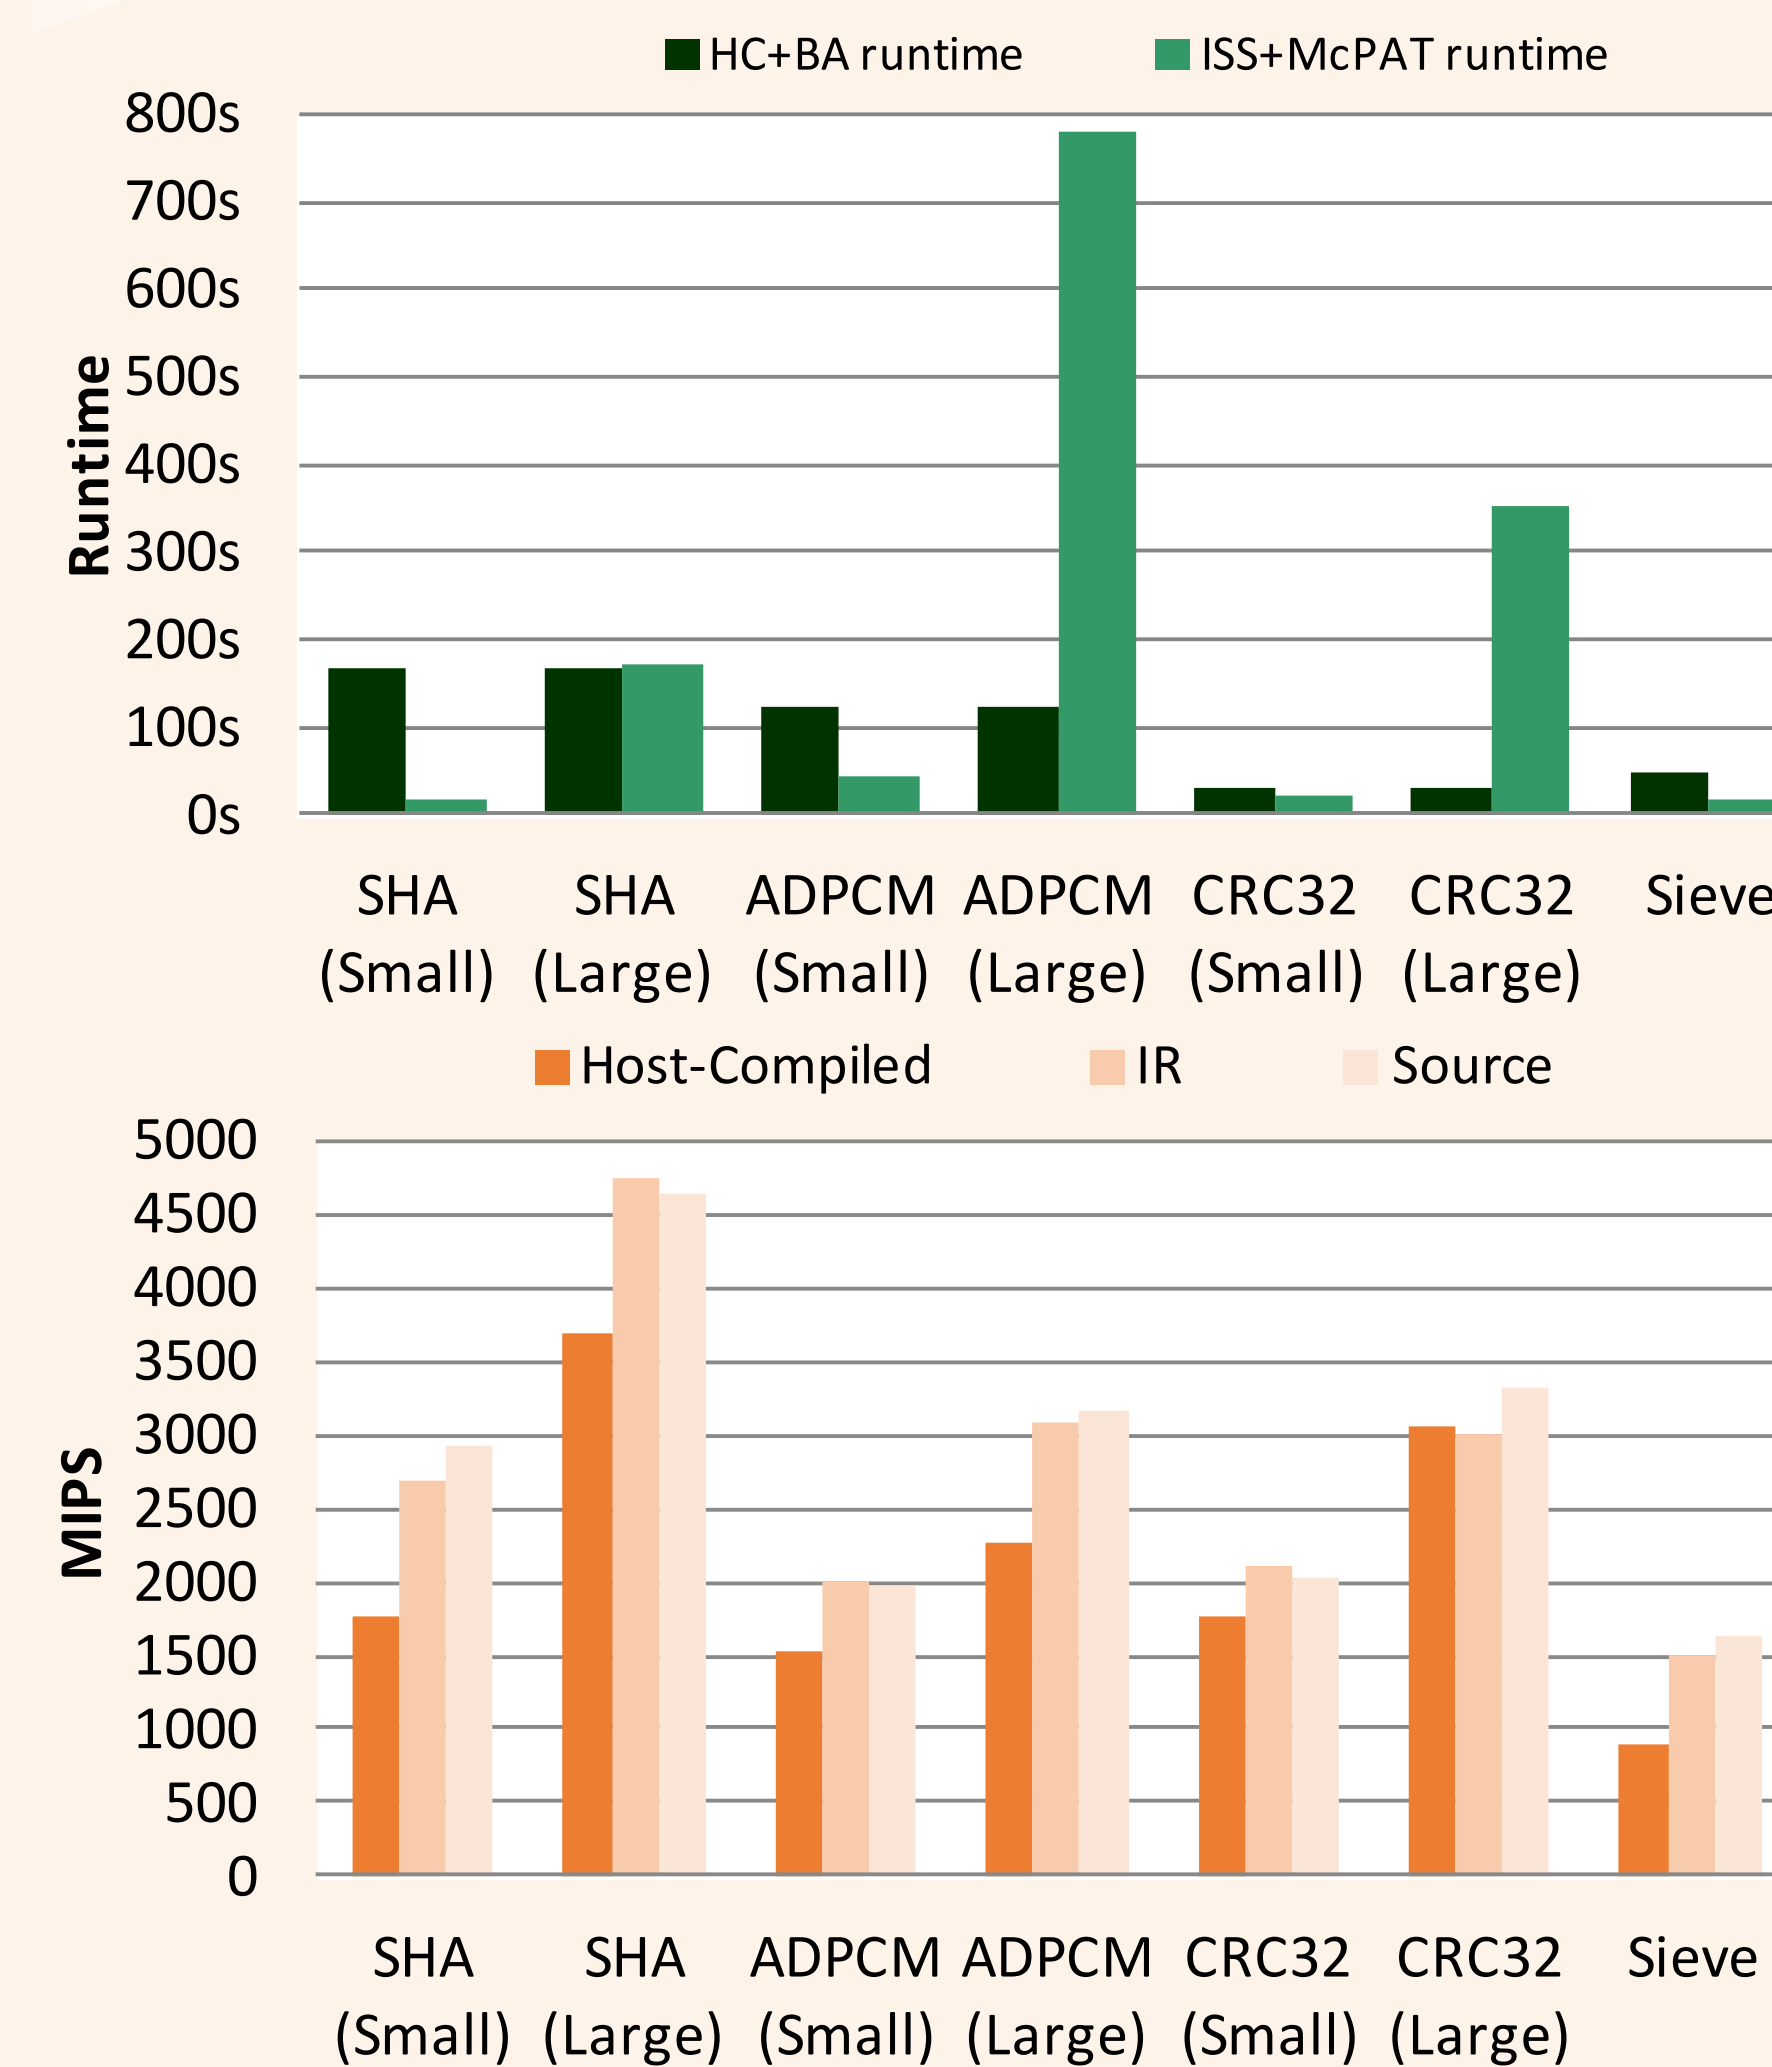
- Modeling above ISS level
 - Compile and execute application natively
 - Annotate application with target timing and power
 - Wrap with SystemC code for platform integration
- Fast and accurate simulation to complement ISS



Experimental Results


Speed Results

- Telecom & security applications [MiBench]
 - SHA, ADPCM, CRC32 & custom Eratosthenes' Sieve
 - Small and large data sets, 10 to 700 million instr.
- One-time back-annotation
 - 3min. to 3s BA runtime
- Host-compiled simulation vs. traditional ISS
 - 2000 MIPS vs. 0.8-1 MIPS
 - Close to source-level speeds



Accuracy Results

- Single- (z4-like) and dual-issue (z6-like) e200 PowerPC
 - No cache, static branch prediction
- Compare against cycle-accurate reference ISS+McPAT
 - >99% average timing and energy accuracy @ 2000 MIPS

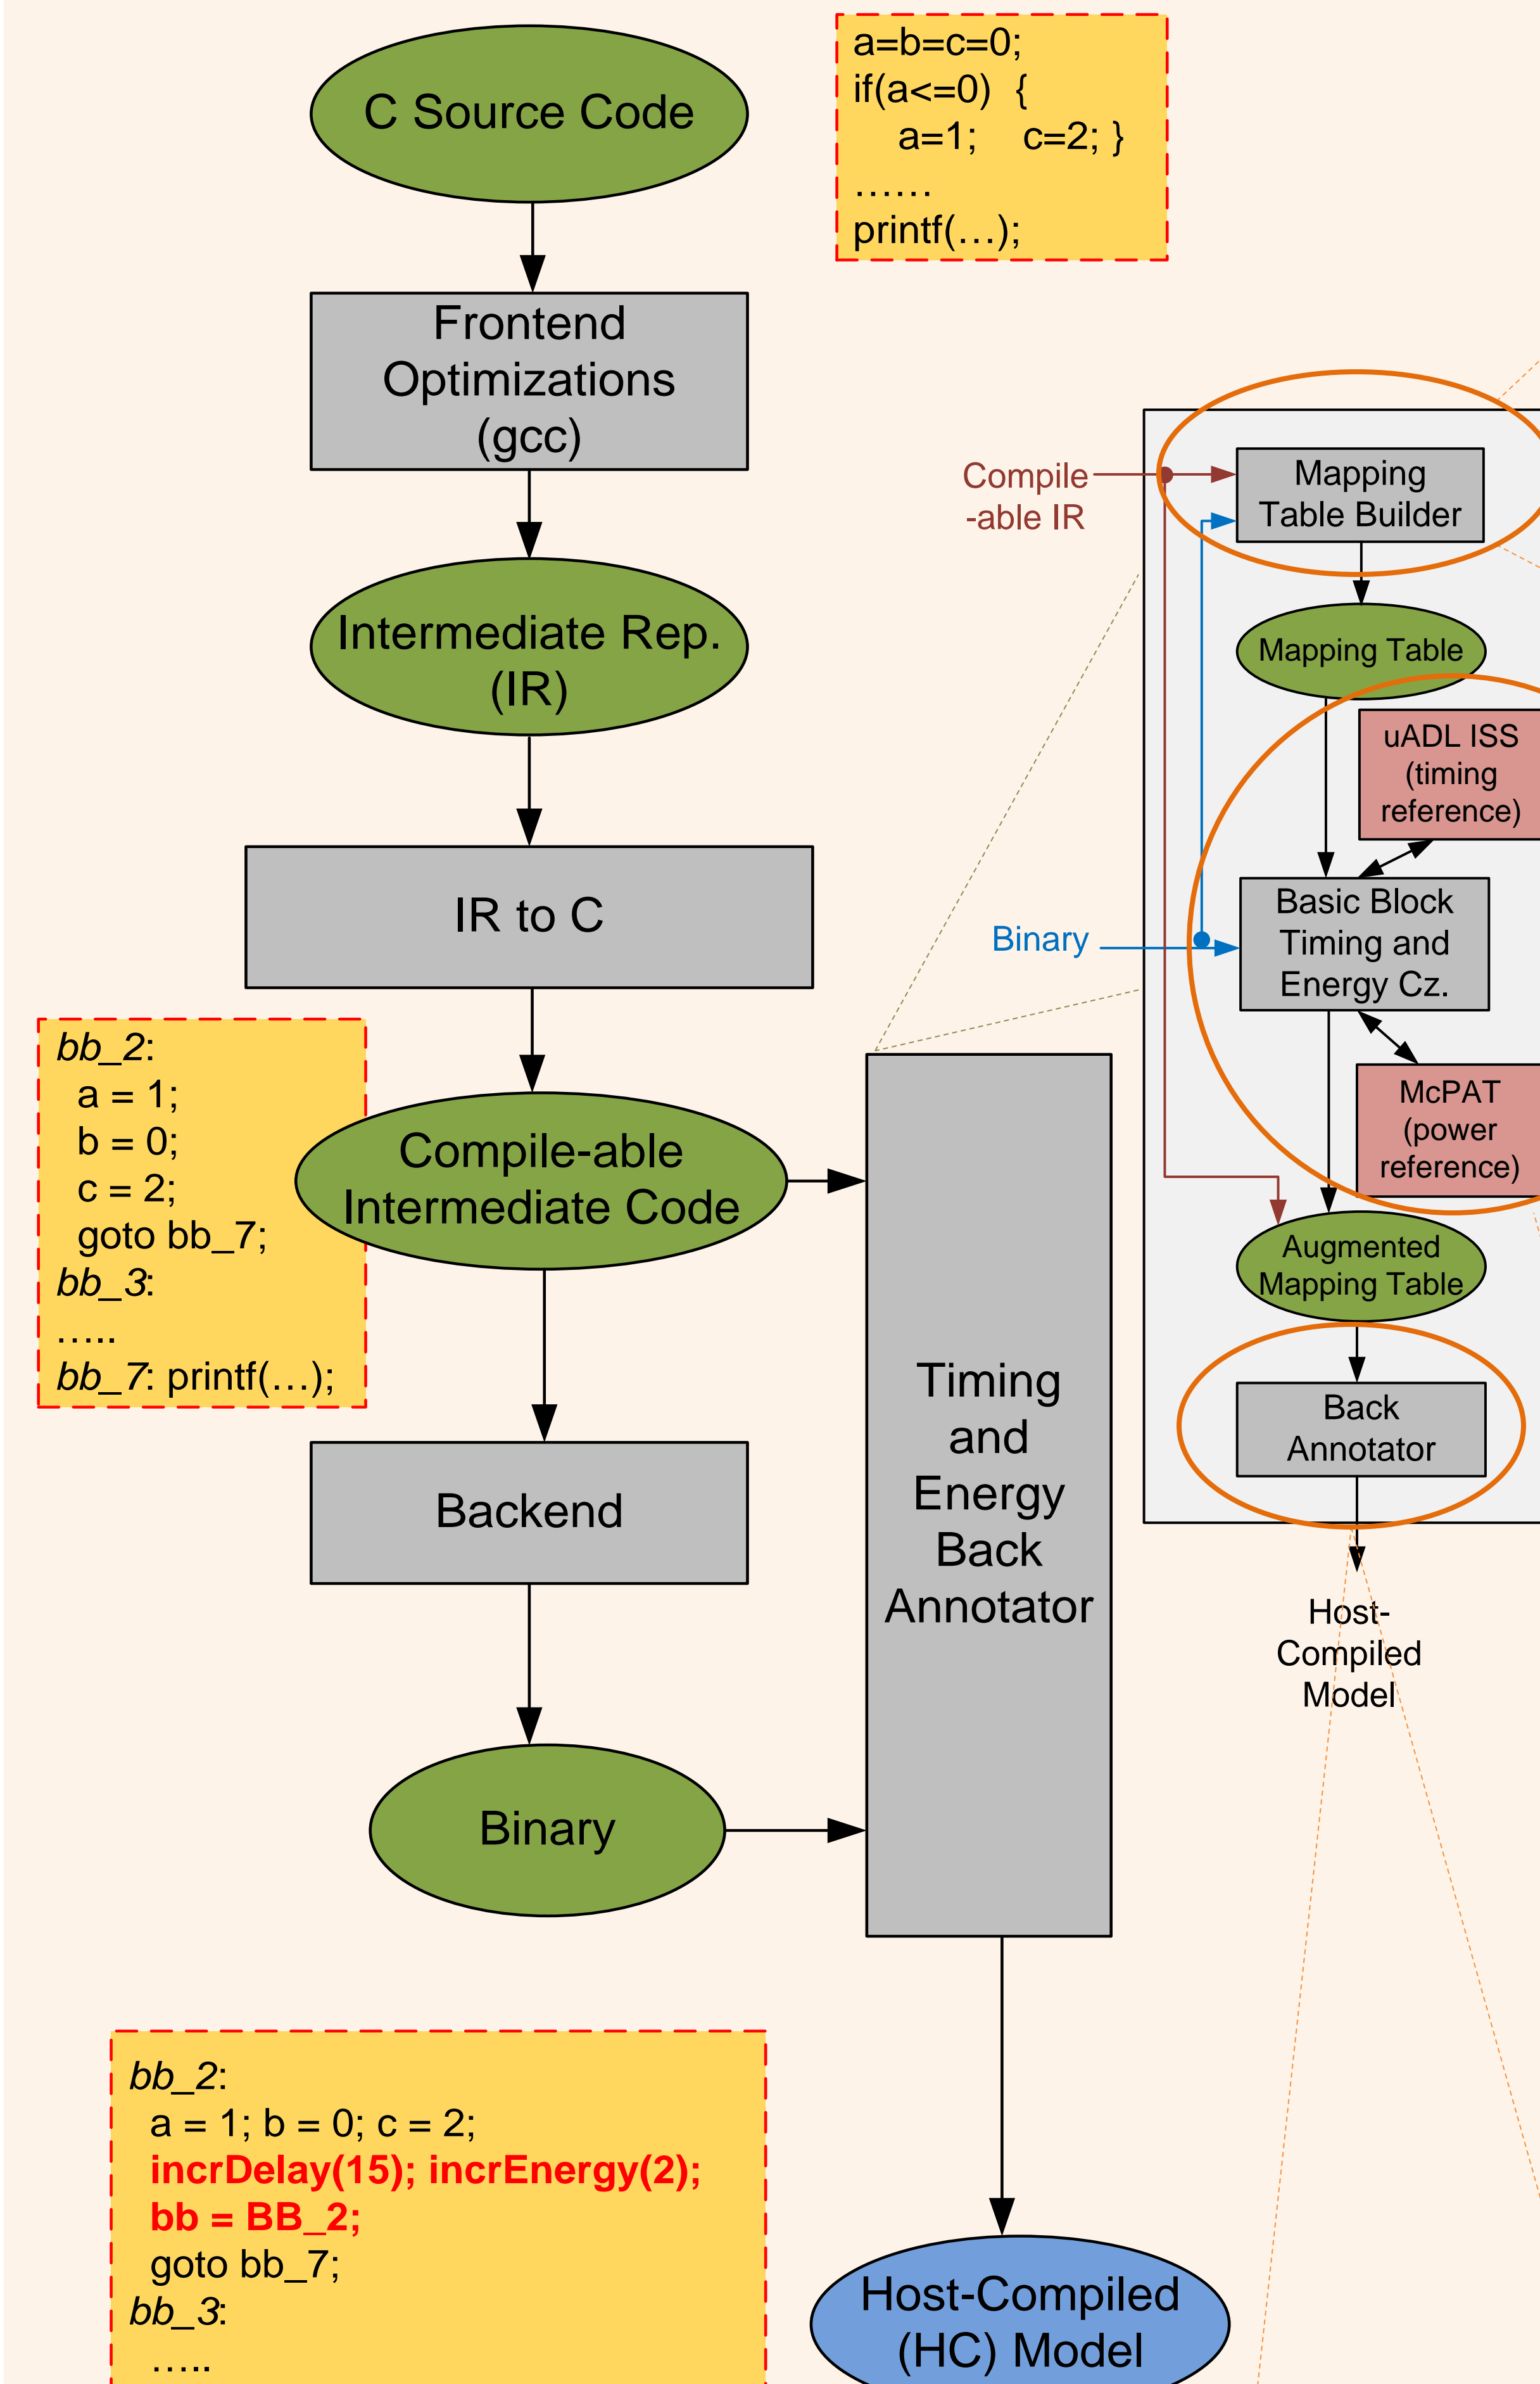


Summary & Conclusions


- Retargetable power/performance back-annotation
- Running at source level speed with >98% accuracy
- Future work
 - Integrated other metrics into host-compiled simulation (thermal, reliability)

Retargetable Back Annotator

Retargetable Back-Annotation Flow



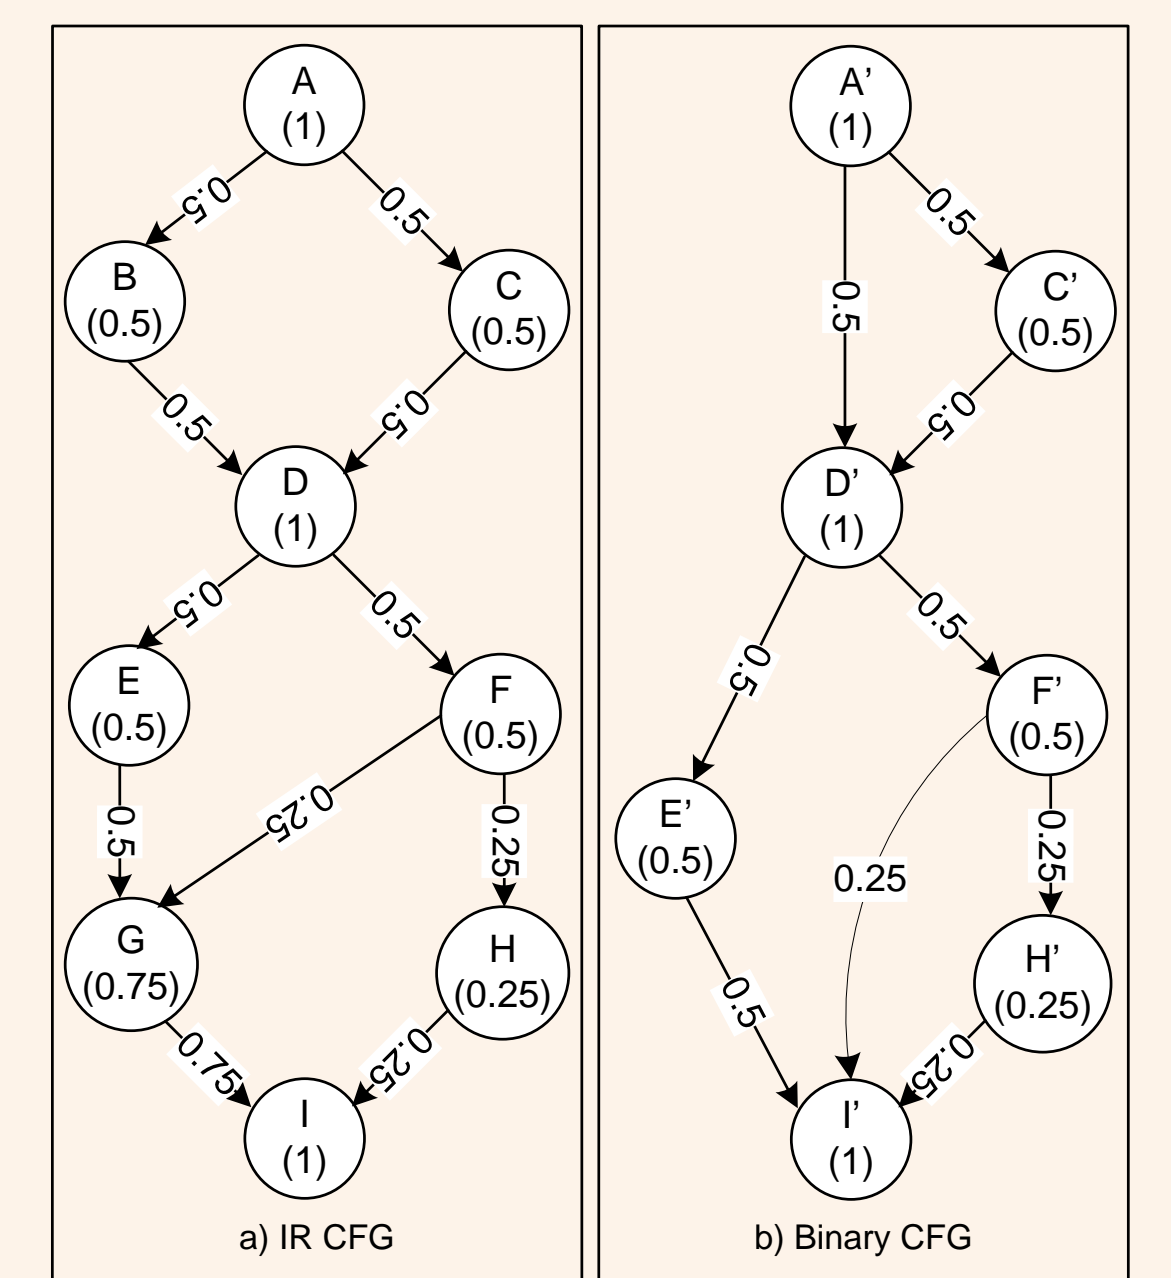
Binary-to-IR Mapping



- Backend optimizations
 - Instruction scheduling
 - Blocks added/removed
 - Predicated execution
 - Control flow mismatches

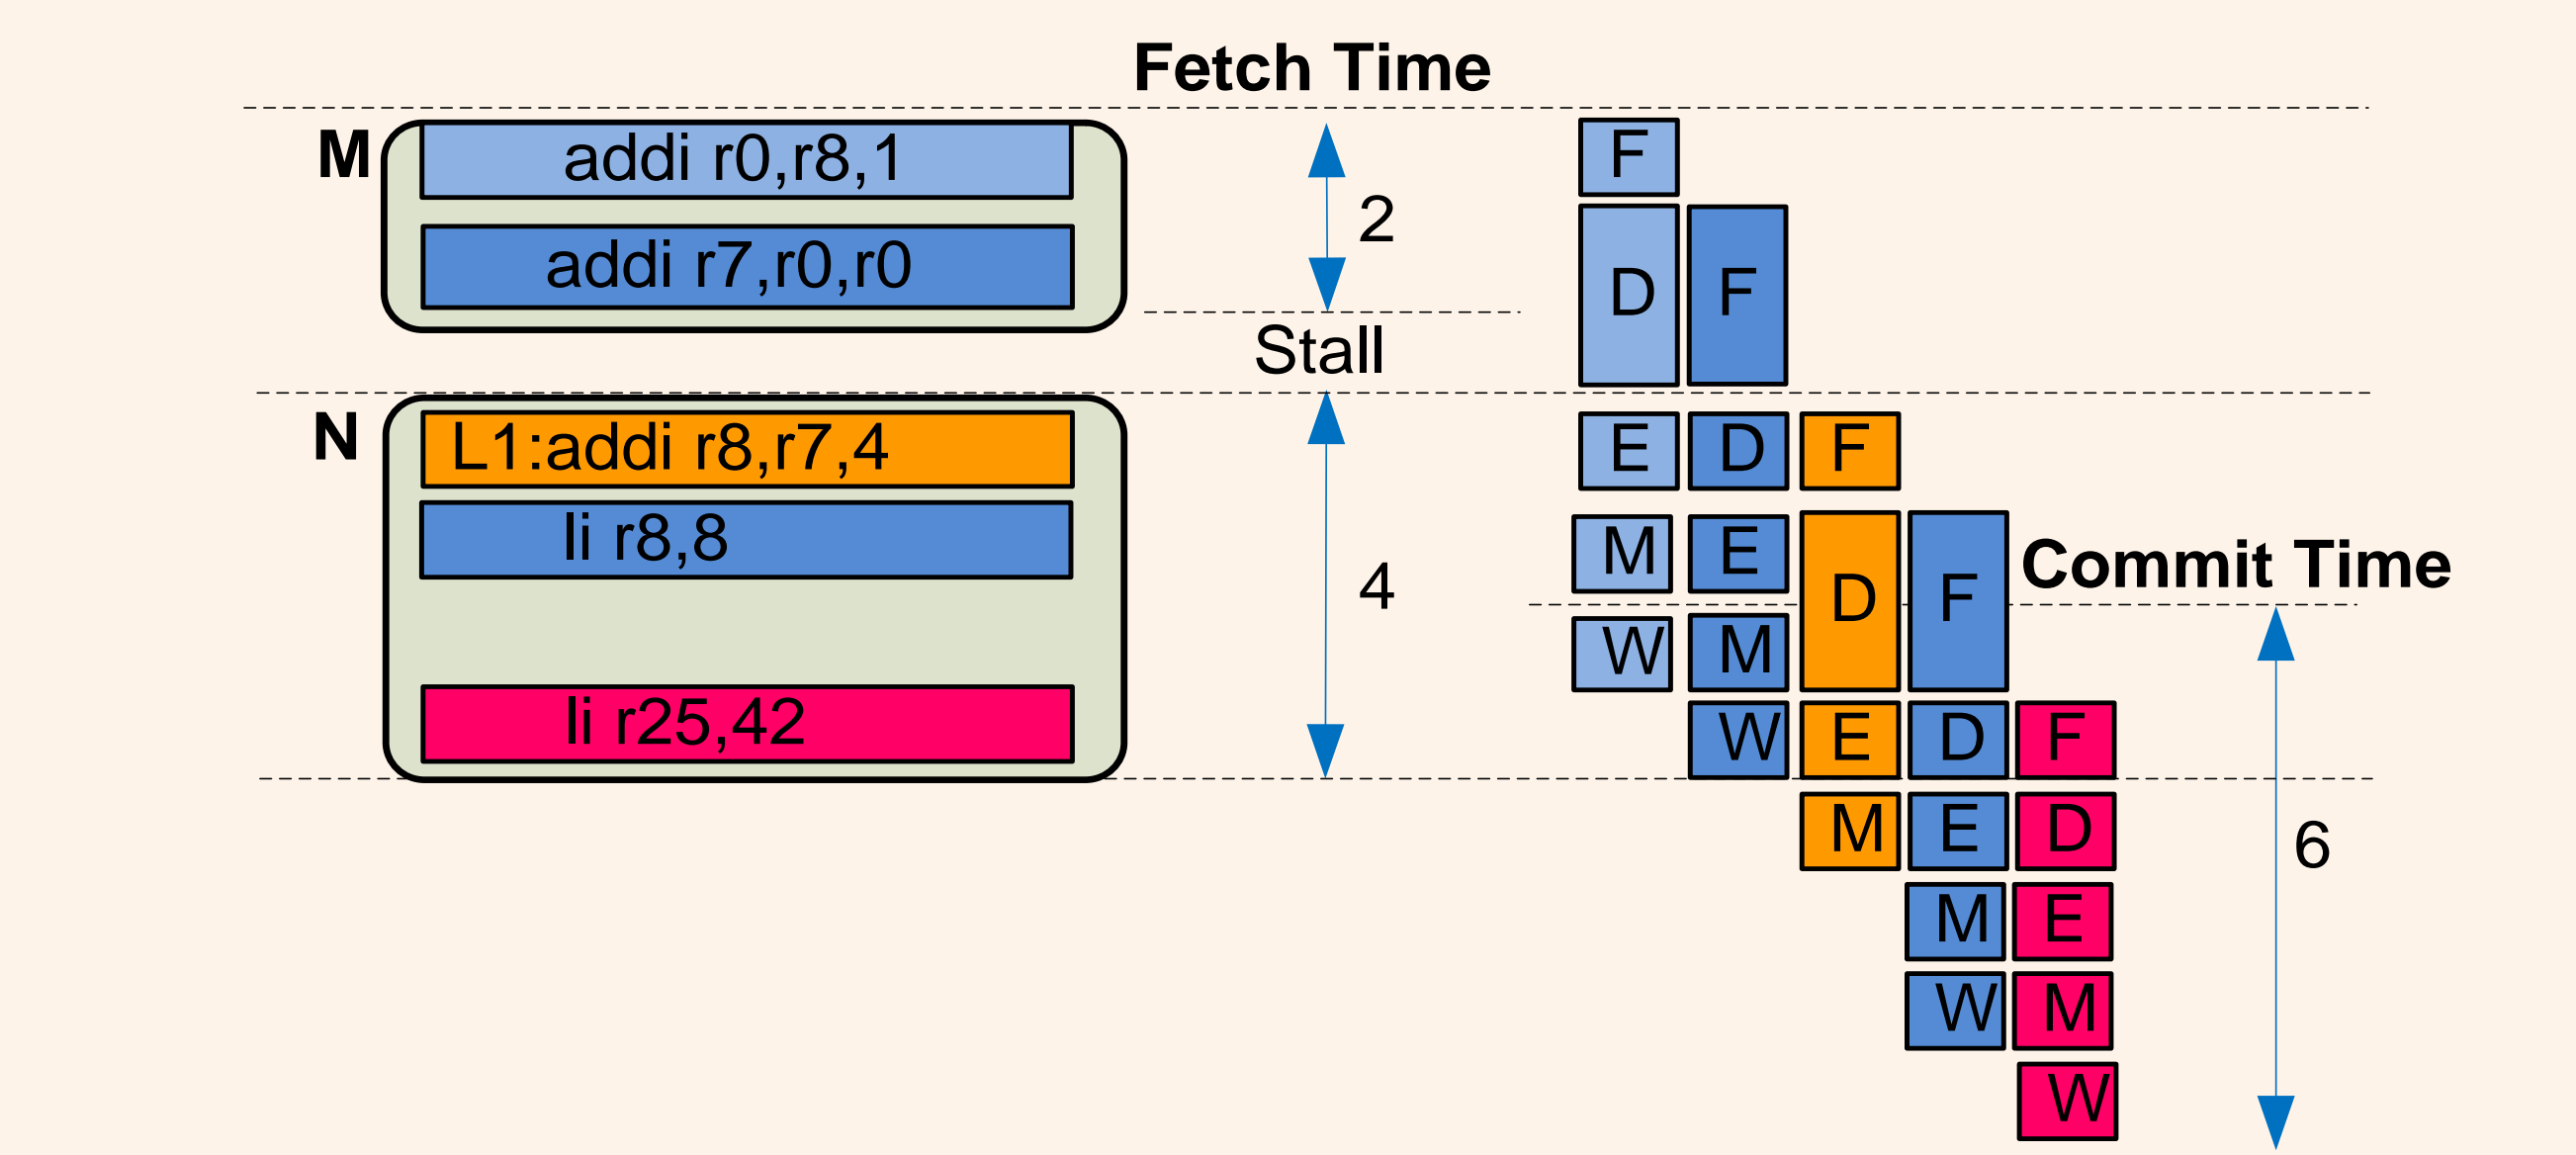
Graph matching heuristic

- Synchronized, recursive depth-first traversal
 - Compatibility: loop and branch nesting levels
 - Cost: sum of unmatched nodes in subgraphs rooted at node
 - Return least-cost mapping between all successors (incl. skips)
- Resolve ambiguities using debug information



Basic Block Characterization

- Metrics depend on system state before BB entry
 - BB1, BB2, BB3
 - SS = A, SS = B, SS = Sys State (registers, mem, pipeline)
- Real execution approximated
 - Pairwise characterization
 - Inter-block stall or overlap is reflected in characterized block
- Function call characterization



Back-annotation into IR

- Path dependent metrics
 - Capture static branch prediction

Technology Transfer

- Reports and publications
 - S. Chakravarty, A. Gerstlauer. "Performance and Power Modeling Case Study", Report, Jan. 2011.
 - S. Chakravarty, A. Gerstlauer. "ADL Driven Back-Annotation and Simulation Methodology for Host Compiled Performance and Power Modeling", Report, July 2011.
 - S. Chakravarty, A. Gerstlauer. "Host Compiled Performance and Power Modeling of Embedded Systems," SRC TECHCON, August 2011.
 - S. Chakravarty, A. Gerstlauer, "Automated, ADL-Driven Generation of Host-Compiled Platform Models," Report, December 2012.
 - S. Chakravarty, Z. Zhao, A. Gerstlauer, "Automated, Retargetable Back-Annotation for Host-Compiled Power and Performance Modeling," CODES+ISSS, Montreal, Canada, September 2013.
- Industrial interactions
 - Leveraging Freescale's open-source ADL technology
 - <http://opensource.freescale.com/fsl-oss-projects/>
 - Regular (bi-weekly) project discussions (Freescale, TI)
 - Internships at Freescale (6/2013-8/2013)
 - Automated discovery of ADL based cycle-accurate reference models